

Data Analysis using spark

Importing and loading Data

```
#df = spark.read.format("mongo").load()
df = spark.read.format("mongo").option("uri",
"mongodb://54.197.41.201/final-project.twitter-collection").load()

df2 = spark.read.format("mongo").option("uri",
"mongodb://54.197.41.201/final-project.youtube-collection").load()
```

Checking if the data is loaded properly

```
print(df.count())
print(df2.count())
```

```
158557
45
```

Twitter sentiment analysis in graphical form

```
df.show()
display(df.select("subjectivity", "polarity"))
```



Showing sample based on the first 1000 rows.



Binning continous variable polarity into categories

```

from pyspark.sql.functions import udf
from pyspark.sql.types import *
def categorizer(polarity):
    if polarity < -0.6:
        return "Very negative"
    elif polarity < -0.1:
        return "Negative"
    elif polarity < 0.1:
        return "Neutral"
    elif polarity < 0.6:
        return "Positive"
    else:
        return "Very Positive"

```

```

bucket_udf = udf(categorizer, StringType() )
bucketed = df.withColumn("polarity_category", bucket_udf("polarity"))
bucketed.show()

```

```


+-----+-----+-----+-----+-----+
+
|          _id|      created_at|      location|      polarity|
screen_name|      subjectivity|      text|      tweet_id|polarity_category|
+-----+-----+-----+-----+-----+
+
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      Argentina|      0.0|
SillviaBak|      0.0|      A este paso no va...|1250191033234849792|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      null|      0.0|
villaverdeh|      0.0|      RT @CiroGomezL: #...|1250191033704632321|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      Philippines|      0.35|
skipDLC|      0.65|      RT @iamkarendavil...|1250191033624748032|      Positive|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|東京都港区.竹島は日本国領土です。|      0.0|
nippon1_|      0.0|RT @nikkei: 中小企業の...|1250191034144952322|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      DM04|      0.0|
ai6yrham|      0.0|      #Hawaii|1250191034241318913|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      null|      0.3|
madamshome|      0.9|      RT @marxdeane: Br...|1250191034258259969|      Positive|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      San Borja, Lima, ...|      0.0|
DavidPbPrz10|      0.0|      RT @sigridbazan: ...|1250191034446802944|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      Sri Lanka|      0.0|
sumanebot|      0.0|      RT @easwaranrutna...|1250191034539220993|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      Permeating|-0.07500000000000001|
NZSilentRage|      0.4|      Empty hospitals
...|1250191034631389184|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      Ankara, Türkiye|      0.0|
kilincbahadir|      0.0|      RT @CemilSahince:...|1250191034698682378|      Neutral|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      null|      -0.6999999999999999|
PleaseThink1776|0.6666666666666666|      RT @PrincessBrava...|1250191035122081794|      Very negativ
e|
|[5e963ac51e0ecc6d...|2020-04-14 22:35:44|      null|      0.0|
nahonak|      0.0|      RT @DrMartyufml: ...|1250191035243905026|      Neutral|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:44|      null|-0.4833333333333333|
okaymarnie|0.7166666666666667|      RT @tanyaboozary:...|1250191035277459463|      Negative|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:44|      Scotland|      0.4|
hoppinghaggis|      1.0|      RT @JamieWoodhous...|1250191035294154760|      Positive|

```

```
|[5e963ac61e0ecc6d...|2020-04-14 22:35:44|Sunny Southern ...|0.0|
SDSUgrad1983|0.0|Hey"Dude" #Hydrox...|1250191035113697280|Neutral|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:44|null|0.050000000000000002|
fernandovd21|0.5|RT @tumbaburross:...|1250191035688349697|Neutral|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:45|味噌ラーメン不毛の地|0.0|
gomadanngou|0.0|RT @sukimaweb: WH...|1250191035893874689|Neutral|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:45|Egypt|0.1875|
YElkabaniBPharm|0.1875|RT @mgranovetter:...|1250191036019806216|Positiv
e|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:45|null|-0.16666666666666666|
BobBack15|0.56666666666666668|RT @monica_sassy:...|1250191035965333504|Negative|
|[5e963ac61e0ecc6d...|2020-04-14 22:35:45|null|-0.8|
SandDollar04|1.0|RT @SandraBundy: ...|1250191036036583424|Very negative|
+-----+-----+-----+-----+
+
only showing top 20 rows
```

Subjectivity score for each polarity category

```
display(bucketed.select("polarity_category","subjectivity"))
```




Aggregated (by avg) in the backend.



Checking variation in polarity overtime

```
display(bucketed.select("created_at","polarity"))
```



Aggregated (by avg) in the backend.
Showing the first 1000 rows.




Casting the type to float (str to float)

```
yt_df = df2.withColumn("viewCount",df2["viewCount"].cast('float'))
yt_df = yt_df.withColumn("likeCount",yt_df["likeCount"].cast('float'))
yt_df = yt_df.withColumn("dislikeCount",yt_df["dislikeCount"].cast('float'))
yt_df = yt_df.withColumn("commentCount",yt_df["commentCount"].cast('float'))
```


Average of view count based on channel title

```
yt_grouped = yt_df.select("viewCount","channelTitle").groupBy("channelTitle").avg().limit(3)
display(yt_grouped)
```



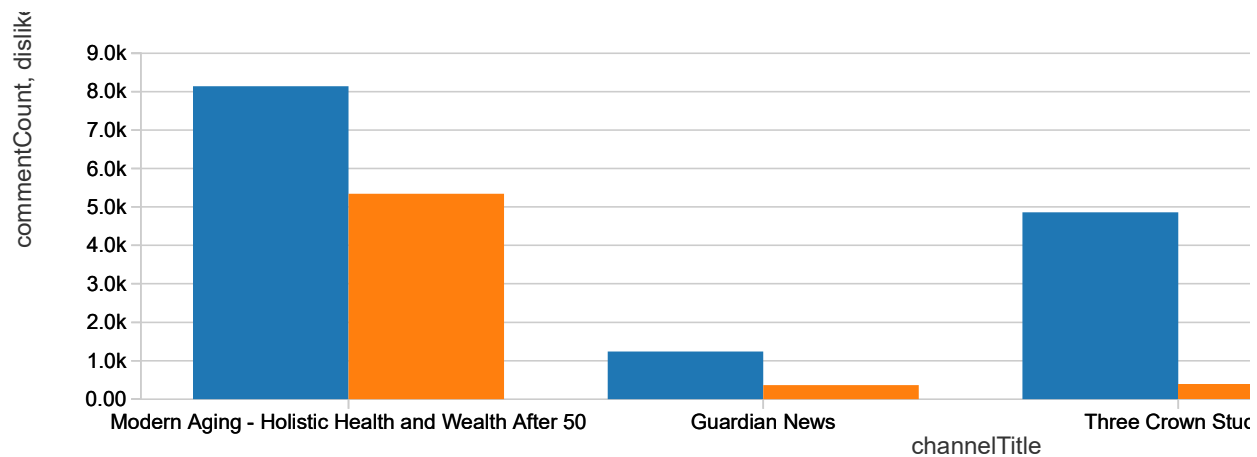
Comparing like count, dislike count, view count and comment count

```
yt_ld_ratio = yt_df.select("likeCount", "dislikeCount", "commentCount", "viewCount")
display(yt_ld_ratio)
```



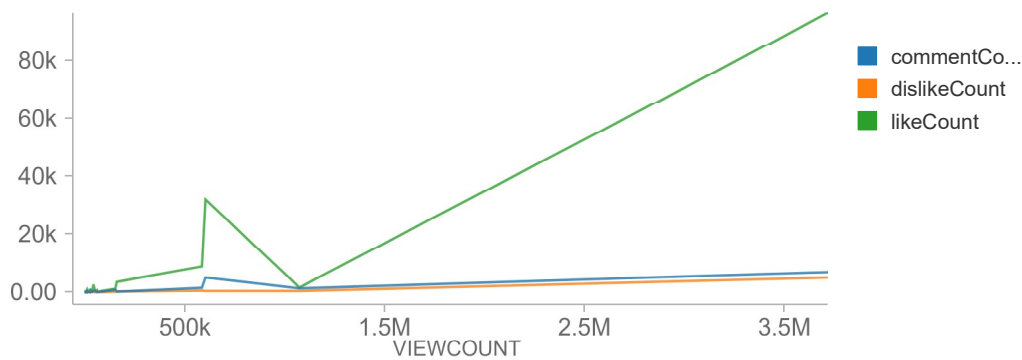
Comparing dislike count and comment count for youtube channels

```
views_likes_comments = yt_df.select("likeCount", "viewCount",
  "commentCount", "dislikeCount", "channelTitle")
vlc_limited = views_likes_comments.limit(5)
display(vlc_limited)
```



Checking the growth of like, dislikes and comments with res

```
display(views_likes_comments)
```



Applying linear regression to predict "Likes"

Coefficients: [0.0143575858110491, 5.49635697215436]

Intercept: -1084.6353744766473

RMSE: 535.135422

r2: 0.991494

prediction	likeCount	features
265.77877294633345	499.0	[4377.0, 53.0]
-7.3924354398246805	51.0	[5616.0, 17.0]
-104.1540122977456	60.0	[5762.0, 4.0]
571.276383110188	49.0	[6224.0, 96.0]
648.9613042216656	278.0	[6771.0, 107.0]

only showing top 5 rows

R Squared (R2) on test data = 0.420769